

Name: **GANESH SANJAY PANDHRE**

Roll Nos : **41** | Div: **D20B**

Aim: To build a Cognitive based application to acquire knowledge through images for a Customer service application/ Insurance/ Healthcare Application/ Smarter Cities/Government etc.

✓ Experiment No: 3

Theory:

Introduction

Cognitive computing refers to systems that can **mimic human intelligence** by using Artificial Intelligence (AI), Machine Learning (ML), and Computer Vision. In this experiment, we focus on **building a cognitive-based application that learns from images** and provides useful knowledge. Such applications are used in **customer service, insurance, healthcare, smarter cities, and government sectors**.

In our case, we continue with the topic from the previous experiment — **Apartment Assistant** — and extend its features by adding **vehicle license plate detection**. This makes the system smarter and more useful for apartment security and management.

Cognitive Applications and Their Role

A **cognitive application** is an AI system that can:

1. **Understand** data (images, text, speech, etc.).
2. **Reason** based on the patterns it finds.
3. **Learn** continuously from new data.
4. **Assist** in decision-making.

In this experiment, the input is **images**, and the system tries to extract **knowledge** from them. Example use cases:

- Recognizing **visitors** entering the apartment.
 - Detecting and reading **vehicle number plates** for parking and security.
 - Extracting **text from bills or notices** using OCR.
-

Concepts Used in This Experiment

1. Image Classification (Computer Vision)

- Deep learning models like **ResNet50** are used to identify objects from images.
- The model recognizes whether the image contains a person, a car, or other objects.

2. Optical Character Recognition (OCR)

- OCR is used to extract text from images (bills, notices, or even vehicle number plates).
- It converts printed or handwritten text into digital form for easy processing.

3. Vehicle License Plate Detection

- A key new addition in this experiment is the **ability to detect and read vehicle plates**.
- This is important for apartments as it can **automate vehicle entry and parking management**.
- The system captures a car's image, detects the plate, and extracts the alphanumeric characters.
- This can be used for **security verification, visitor vehicle logging, and automatic gate access**.

4. Knowledge Acquisition from Images

- Beyond detection, the system interprets the extracted data to make it useful.
- Example: If the license plate number is recognized, it can be checked against the apartment's database of registered vehicles.

Applications in Apartment Assistant

By extending the Apartment Assistant with **vehicle plate detection**, the system becomes more practical:

- **Visitor Identification** → Security can identify and log visitors.
 - **Vehicle Management** → Detect cars entering, check if they are registered, and allow/deny access automatically.
 - **Billing and Notices** → OCR extracts details from electricity bills, maintenance slips, or apartment notices.
-

Working of the Experiment

1. **Input** → User uploads an image (visitor photo, car image, or bill).
 2. **Processing** →
 - Deep learning model identifies objects.
 - OCR extracts text (from number plates or documents).
 3. **Output** →
 - Recognized object category (Car, Person, Document).
 - Extracted license plate text or bill details.
 - Useful knowledge like identifying whether a vehicle is registered or not.
-

Advantages of Apartment Assistant with Vehicle Plate Detection

- Provides **better security** by tracking vehicles entering the apartment.
 - Saves **manual effort** of security guards for logging visitor cars.
 - Helps in **automation of parking management**.
 - Can be combined with other features (face recognition, bill scanning) to form a **complete smart apartment solution**.
-

```
import os
import cv2
import shutil
import numpy as np
import pandas as pd
from glob import glob
import matplotlib.pyplot as plt
import xml.etree.ElementTree as xet
from sklearn.model_selection import train_test_split
```

```
import os
import cv2
import shutil
import numpy as np
import pandas as pd
from glob import glob
import matplotlib.pyplot as plt
import xml.etree.ElementTree as xet
from sklearn.model_selection import train_test_split
```

```
import torch
```

```
print(f'{torch.cuda.is_available() = }')
print(f'{torch.cuda.device_count() = }')
```

```

torch.cuda.is_available() = True
torch.cuda.device_count() = 1

```

```
!pip install ultralytics
```

 [Show hidden output](#)


```
!pip install -U ipywidgets
```

 [Show hidden output](#)

```

from google.colab import files
files.upload()

```

 kaggle.json

- **kaggle.json**(application/json) - 67 bytes, last modified: 8/30/2025 - 100% done

Saving kaggle.json to kaggle.json


```
{'kaggle.json': b'{"username": "robstark143", "key": "cfc32d778a754be45dee9d7207d5eca4"}'}
```

```

!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json

```

```
!kaggle datasets download -d andrewmvd/car-plate-detection
```

 Dataset URL: <https://www.kaggle.com/datasets/andrewmvd/car-plate-detection>
 License(s): CC0-1.0
 Downloading car-plate-detection.zip to /content
 57% 116M/203M [00:00<00:00, 1.18GB/s]
 100% 203M/203M [00:00<00:00, 800MB/s]

```
!unzip car-plate-detection.zip -d car-plate-detection
```

 [Show hidden output](#)

```
dataset_path = '/content/car-plate-detection'
```

```
import re
```

```

def the_number_in_the_string(filename):
    """
    Extracts the first sequence of digits from the given filename string and returns it as an integer.
    If no digits are found, returns 0.

    Parameters:
    filename (str): The input string to search for digits.

    Returns:
    int: The first sequence of digits found in the input string, or 0 if no digits are found.
    """
    # Search for the first occurrence of one or more digits in the filename
    match = re.search(r'(\d+)', filename)

    # If a match is found, return the matched number as an integer
    if match:
        return int(match.group(0))
    # If no match is found, return 0
    else:
        return 0

# Example usage

```

```
print(the_number_in_the_string("file123.txt")) # Output: 123
print(the_number_in_the_string("no_numbers_here")) # Output: 0
```

```
→ 123
   0
```

```
import os
import cv2
import pandas as pd
import xml.etree.ElementTree as xet
from glob import glob

# Initialize a dictionary to store labels and image information
labels_dict = dict(
    img_path=[],
    xmin=[],
    xmax=[],
    ymin=[],
    ymax=[],
    img_w=[],
    img_h=[]
)

# Get the list of XML files from the annotations directory
xml_files = glob(f'{dataset_path}/annotations/*.xml')

# Process each XML file, sorted by the numerical value in the filename
for filename in sorted(xml_files, key=the_number_in_the_string):
    # Parse the XML file
    info = xet.parse(filename)
    root = info.getroot()

    # Find the 'object' element in the XML and extract bounding box information
    member_object = root.find('object')
    labels_info = member_object.find('bndbox')
    xmin = int(labels_info.find('xmin').text)
    xmax = int(labels_info.find('xmax').text)
    ymin = int(labels_info.find('ymin').text)
    ymax = int(labels_info.find('ymax').text)

    # Get the image filename and construct the full path to the image
    img_name = root.find('filename').text
    img_path = os.path.join(dataset_path, 'images', img_name)

    # Append the extracted information to the respective lists in the dictionary
    labels_dict['img_path'].append(img_path)
    labels_dict['xmin'].append(xmin)
    labels_dict['xmax'].append(xmax)
    labels_dict['ymin'].append(ymin)
    labels_dict['ymax'].append(ymax)

    # Read the image to get its dimensions
    height, width, _ = cv2.imread(img_path).shape
    labels_dict['img_w'].append(width)
    labels_dict['img_h'].append(height)

# Convert the dictionary to a pandas DataFrame
alldata = pd.DataFrame(labels_dict)

# Display the DataFrame
alldata
```



	img_path	xmin	xmax	ymin	ymax	img_w	img_h	
0	/content/car-plate-detection/images/Cars0.png	226	419	125	173	500	268	
1	/content/car-plate-detection/images/Cars1.png	134	262	128	160	400	248	
2	/content/car-plate-detection/images/Cars2.png	229	270	176	193	400	400	
3	/content/car-plate-detection/images/Cars3.png	142	261	128	157	400	225	
4	/content/car-plate-detection/images/Cars4.png	156	503	82	253	590	350	
...	
428	/content/car-plate-detection/images/Cars428.png	142	258	128	157	400	225	
429	/content/car-plate-detection/images/Cars429.png	86	208	166	195	301	400	
430	/content/car-plate-detection/images/Cars430.png	38	116	159	197	400	225	
431	/content/car-plate-detection/images/Cars431.png	55	343	82	147	400	192	
432	/content/car-plate-detection/images/Cars432.png	95	196	258	284	467	300	

433 rows × 7 columns

Next steps:

[Generate code with alldata](#)[View recommended plots](#)[New interactive sheet](#)

```

from sklearn.model_selection import train_test_split

# Split the data into training and test sets
# Use 10% of the data for the test set
train, test = train_test_split(alldata, test_size=1/10, random_state=42)

# Split the training data further into training and validation sets
# Use 8/9 of the remaining data for the training set, resulting in an 80/10/10 split overall
train, val = train_test_split(train, train_size=8/9, random_state=42)

# Print the number of samples in each set
print(f'''
    len(train) = {len(train)}
    len(val) = {len(val)}
    len(test) = {len(test)}
''')
```



```

len(train) = 345
len(val) = 44
len(test) = 44
```

```

import os
import shutil
import pandas as pd

# Remove the 'datasets' directory if it exists
if os.path.exists('datasets'):
    shutil.rmtree('datasets')
```

```

def make_split_folder_in_yolo_format(split_name, split_df):
    """
    Creates a folder structure for a dataset split (train/val/test) in YOLO format.

    Parameters:
    split_name (str): The name of the split (e.g., 'train', 'val', 'test').
    split_df (pd.DataFrame): The DataFrame containing the data for the split.

    The function will create 'labels' and 'images' subdirectories under 'datasets/cars_license_plate/{split_name}',
```

```

and save the corresponding labels and images in YOLO format.
"""

labels_path = os.path.join('datasets', 'cars_license_plate_new', split_name, 'labels')
images_path = os.path.join('datasets', 'cars_license_plate_new', split_name, 'images')

# Create directories for labels and images
os.makedirs(labels_path)
os.makedirs(images_path)

# Iterate over each row in the DataFrame
for _, row in split_df.iterrows():
    img_name, img_extension = os.path.splitext(os.path.basename(row['img_path']))

    # Calculate YOLO format bounding box coordinates
    x_center = (row['xmin'] + row['xmax']) / 2 / row['img_w']
    y_center = (row['ymin'] + row['ymax']) / 2 / row['img_h']
    width = (row['xmax'] - row['xmin']) / row['img_w']
    height = (row['ymax'] - row['ymin']) / row['img_h']

    # Save the label in YOLO format
    label_path = os.path.join(labels_path, f'{img_name}.txt')
    with open(label_path, 'w') as file:
        file.write(f"{x_center:.4f} {y_center:.4f} {width:.4f} {height:.4f}\n")

    # Copy the image to the images directory
    shutil.copy(row['img_path'], os.path.join(images_path, img_name + img_extension))

print(f"Created '{images_path}' and '{labels_path}'")

```

```

# Create YOLO format folders for train, validation, and test splits
make_split_folder_in_yolo_format("train", train)
make_split_folder_in_yolo_format("val", val)
make_split_folder_in_yolo_format("test", test)

```



Created 'datasets/cars_license_plate_new/train/images' and 'datasets/cars_license_plate_new/train/labels'
 Created 'datasets/cars_license_plate_new/val/images' and 'datasets/cars_license_plate_new/val/labels'
 Created 'datasets/cars_license_plate_new/test/images' and 'datasets/cars_license_plate_new/test/labels'

```
os.getcwd()
```



'/content'

```

import os
import cv2
import matplotlib.pyplot as plt

# Directory paths
image_dir = 'datasets/cars_license_plate_new/train/images'
label_dir = 'datasets/cars_license_plate_new/train/labels'

# Get the first image file
image_files = sorted(os.listdir(image_dir))
first_image_file = image_files[0]

# Construct paths for the image and its corresponding label
image_path = os.path.join(image_dir, first_image_file)
label_path = os.path.join(label_dir, os.path.splitext(first_image_file)[0] + '.txt')

# Load the image using OpenCV
image = cv2.imread(image_path)
# Convert the image from BGR (OpenCV default) to RGB (matplotlib default)
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

# Read the label file to get bounding box information
with open(label_path, 'r') as f:

```

```

lines = f.readlines()

# Plot the bounding box on the image
for line in lines:
    # Parse the label file line to extract bounding box information
    class_id, x_center, y_center, width, height = map(float, line.strip().split())
    img_height, img_width, _ = image.shape

    # Convert YOLO format to bounding box format
    x_center *= img_width
    y_center *= img_height
    width *= img_width
    height *= img_height

    # Calculate the top-left and bottom-right coordinates of the bounding box
    x1 = int(x_center - width / 2)
    y1 = int(y_center - height / 2)
    x2 = int(x_center + width / 2)
    y2 = int(y_center + height / 2)

    # Draw the bounding box on the image using a green rectangle
    cv2.rectangle(image, (x1, y1), (x2, y2), (0, 255, 0), 2)

# Display the image with bounding box using matplotlib
plt.imshow(image)
plt.axis('off') # Hide the axis
plt.show() # Display the image

```



```

# Define the content of the datasets.yaml file
datasets_yaml = '''
path: cars_license_plate_new

train: train/images
val: val/images
test: test/images



# number of classes
nc: 1

# class names
names: ['license_plate']
'''



# Write the content to the datasets.yaml file
with open('datasets.yaml', 'w') as file:
    file.write(datasets_yaml)

```

```
from ultralytics import YOLO
model = YOLO('yolov8n.pt')
```

Creating new Ultralytics Settings v0.0.6 file 
 View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.json'
 Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see <https://docs.ultralytics.com>.
 Downloading <https://github.com/ultralytics/assets/releases/download/v8.3.0/yolov8n.pt> to 'yolov8n.pt': 100% 

```
model.train(
    data='datasets.yaml', # Path to the dataset configuration file
    epochs=100,           # Number of training epochs
    batch=16,             # Batch size
    device='cuda',        # Use GPU for training
    imgsz=320,            # Image size (width and height) for training
    cache=True            # Cache images for faster training
)
```

Ultralytics 8.3.189  Python-3.12.11 torch-2.8.0+cu126 CUDA:0 (Tesla T4, 15095MiB)
engine/trainer: agnostic_nms=False, amp=True, augment=False, auto_augment=randaugument, batch=16, bgr=0.0, box=7.5,
 Downloading <https://ultralytics.com/assets/Arial.ttf> to '/root/.config/Ultralytics/Arial.ttf': 100%  75
 Overriding model.yaml nc=80 with nc=1

	from	n	params	module	arguments
0	-1	1	464	ultralytics.nn.modules.conv.Conv	[3, 16, 3, 2]
1	-1	1	4672	ultralytics.nn.modules.conv.Conv	[16, 32, 3, 2]
2	-1	1	7360	ultralytics.nn.modules.block.C2f	[32, 32, 1, True]
3	-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
4	-1	2	49664	ultralytics.nn.modules.block.C2f	[64, 64, 2, True]
5	-1	1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
6	-1	2	197632	ultralytics.nn.modules.block.C2f	[128, 128, 2, True]
7	-1	1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
8	-1	1	460288	ultralytics.nn.modules.block.C2f	[256, 256, 1, True]
9	-1	1	164608	ultralytics.nn.modules.block.SPPF	[256, 256, 5]
10	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
11	[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
12	-1	1	148224	ultralytics.nn.modules.block.C2f	[384, 128, 1]
13	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
14	[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
15	-1	1	37248	ultralytics.nn.modules.block.C2f	[192, 64, 1]
16	-1	1	36992	ultralytics.nn.modules.conv.Conv	[64, 64, 3, 2]
17	[-1, 12]	1	0	ultralytics.nn.modules.conv.Concat	[1]
18	-1	1	123648	ultralytics.nn.modules.block.C2f	[192, 128, 1]
19	-1	1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
20	[-1, 9]	1	0	ultralytics.nn.modules.conv.Concat	[1]
21	-1	1	493056	ultralytics.nn.modules.block.C2f	[384, 256, 1]
22	[15, 18, 21]	1	751507	ultralytics.nn.modules.head.Detect	[1, [64, 128, 256]]

Model summary: 129 layers, 3,011,043 parameters, 3,011,027 gradients, 8.2 GFLOPs

Transferred 319/355 items from pretrained weights

Freezing layer 'model.22.dfl.conv.weight'

AMP: running Automatic Mixed Precision (AMP) checks...


Downloading <https://github.com/ultralytics/assets/releases/download/v8.3.0/yolo11n.pt> to 'yolo11n.pt': 100% 

AMP: checks passed 

train: Fast image access  (ping: 0.0±0.0 ms, read: 2208.3±1283.1 MB/s, size: 607.4 KB)


train: Scanning /content/datasets/cars_license_plate_new/train/labels... 345 images, 0 backgrounds, 0 corrupt: 100%


train: New cache created: /content/datasets/cars_license_plate_new/train/labels.cache

WARNING  cache='ram' may produce non-deterministic training results. Consider cache='disk' as a deterministic alt


train: Caching images (0.1GB RAM): 100%  345/345 58.8it/s 5.9s


augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, method='weig

val: Fast image access  (ping: 0.0±0.0 ms, read: 1154.2±1237.2 MB/s, size: 402.1 KB)

val: Scanning /content/datasets/cars_license_plate_new/val/labels... 44 images, 0 backgrounds, 0 corrupt: 100% 

val: New cache created: /content/datasets/cars_license_plate_new/val/labels.cache

WARNING  cache='ram' may produce non-deterministic training results. Consider cache='disk' as a deterministic alt

val: Caching images (0.0GB RAM): 100%  44/44 72.4it/s 0.6s

Plotting labels to runs/detect/train/labels.jpg...

optimizer: 'optimizer=auto' found, ignoring 'lr=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0'

optimizer: AdamW(lr=0.002, momentum=0.9) with parameter groups 57 weight(decay=0.0), 64 weight(decay=0.0005), 63 bi

Image sizes 320 train, 320 val

Using 2 dataloader workers

Logging results to runs/detect/train

Starting training for 100 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/100	0.594G	1.758	3.007	1.347	11	320: 100% 22/22 4.0it/s 5.4s


```
import os
import pandas as pd
import matplotlib.pyplot as plt
from glob import glob

# Find the most recent training log directory
log_dir = max(glob('runs/detect/train*'), key=lambda x: os.path.getmtime(x))

# Load the training results from the CSV file
results = pd.read_csv(os.path.join(log_dir, 'results.csv'))
results.columns = results.columns.str.strip() # Remove any leading/trailing whitespace from column names

# Extract epochs and accuracy metrics
epochs = results.index + 1 # Epochs are zero-indexed, so add 1
mAP_0_5 = results['metrics/mAP50(B)'] # Mean Average Precision at IoU=0.5
mAP_0_5_0_95 = results['metrics/mAP50-95(B)'] # Mean Average Precision at IoU=0.5:0.95

# Plot the accuracy over epochs
plt.figure(figsize=(10, 5))
plt.plot(epochs, mAP_0_5, label='mAP@0.5')
plt.plot(epochs, mAP_0_5_0_95, label='mAP@0.5:0.95')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Accuracy Over Epochs')
plt.legend()
plt.grid(True)
plt.show()
```



```
# Save the trained model
model.save('best_license_plate_model.pt')
```

```
import cv2
import matplotlib.pyplot as plt
from ultralytics import YOLO
```

```
def predict_and_plot(path_test_car):
    """
    Predicts and plots the bounding boxes on the given test image using the trained YOLO model.

    Parameters:
    path_test_car (str): Path to the test image file.
    """
    # Perform prediction on the test image using the model
    results = model.predict(path_test_car, device='cpu')

    # Load the image using OpenCV
    image = cv2.imread(path_test_car)
    # Convert the image from BGR (OpenCV default) to RGB (matplotlib default)
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

    # Extract the bounding boxes and labels from the results
    for result in results:
        for box in result.bboxes:
            # Get the coordinates of the bounding box
            x1, y1, x2, y2 = map(int, box.xyxy[0])
            # Get the confidence score of the prediction
            confidence = box.conf[0]

            # Draw the bounding box on the image
            cv2.rectangle(image, (x1, y1), (x2, y2), (0, 255, 0), 2)
            # Draw the confidence score near the bounding box
            cv2.putText(image, f'{confidence*100:.2f}%', (x1, y1 - 10),
                        cv2.FONT_HERSHEY_SIMPLEX, 0.9, (255, 0, 0), 2)

    # Plot the image with bounding boxes
    plt.imshow(image)
    plt.axis('off') # Hide the axis
    plt.show() # Display the image
```

```
predict_and_plot(test.iloc[0].img_path)
```




image 1/1 /content/car-plate-detection/images/Cars425.png: 160x320 1 license_plate, 47.8ms
Speed: 0.7ms preprocess, 47.8ms inference, 24.9ms postprocess per image at shape (1, 3, 160, 320)



```
predict_and_plot(test.iloc[15].img_path)
```



image 1/1 /content/car-plate-detection/images/Cars395.png: 256x320 1 license_plate, 70.7ms
Speed: 1.2ms preprocess, 70.7ms inference, 0.9ms postprocess per image at shape (1, 3, 256, 320)



```
predict_and_plot(test.iloc[18].img_path)
```



image 1/1 /content/car-plate-detection/images/Cars9.png: 256x320 1 license_plate, 79.0ms
Speed: 1.4ms preprocess, 79.0ms inference, 0.9ms postprocess per image at shape (1, 3, 256, 320)



```
!pip3 install pytesseract
```



```
Collecting pytesseract
  Downloading pytesseract-0.3.13-py3-none-any.whl.metadata (11 kB)
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.12/dist-packages (from pytesseract) (25.0)
Requirement already satisfied: Pillow>=8.0.0 in /usr/local/lib/python3.12/dist-packages (from pytesseract) (11.3.0)
Downloading pytesseract-0.3.13-py3-none-any.whl (14 kB)
Installing collected packages: pytesseract
Successfully installed pytesseract-0.3.13
```

```
import pytesseract
from pytesseract import Output
```

```

import cv2
import matplotlib.pyplot as plt
from ultralytics import YOLO
import pytesseract
from pytesseract import Output

def predict_and_plot(path_test_car):
    """
    Predicts and plots the bounding boxes on the given test image using the trained YOLO model.
    Also performs OCR on the detected bounding boxes to extract text.

    Parameters:
    path_test_car (str): Path to the test image file.
    """
    # Perform prediction on the test image using the model
    results = model.predict(path_test_car, device='cpu')

    # Load the image using OpenCV
    image = cv2.imread(path_test_car)
    # Convert the image from BGR (OpenCV default) to RGB (matplotlib default)
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

    # Extract the bounding boxes and labels from the results
    for result in results:
        for box in result.bboxes:
            # Get the coordinates of the bounding box
            x1, y1, x2, y2 = map(int, box.xyxy[0])
            # Get the confidence score of the prediction
            confidence = box.conf[0]

            # Draw the bounding box on the image
            cv2.rectangle(image, (x1, y1), (x2, y2), (0, 255, 0), 2)
            # Draw the confidence score near the bounding box
            cv2.putText(image, f'{confidence*100:.2f}%', (x1, y1 - 10),
                        cv2.FONT_HERSHEY_SIMPLEX, 0.9, (255, 0, 0), 2)

            # Crop the bounding box from the image for OCR
            roi = image[y1:y2, x1:x2]

            # Perform OCR on the cropped image
            text = pytesseract.image_to_string(roi, config='--psm 6')
            print(f"Detected text: {text}")

    # Plot the image with bounding boxes
    plt.imshow(image)
    plt.axis('off') # Hide the axis
    plt.show() # Display the image

```

```

predict_and_plot(test.iloc[0].img_path)

```



image 1/1 /content/car-plate-detection/images/Cars425.png: 160x320 1 license_plate, 55.3ms
Speed: 2.7ms preprocess, 55.3ms inference, 1.2ms postprocess per image at shape (1, 3, 160, 320)
Detected text: G526 JHD



```
predict_and_plot(test.iloc[2].img_path)
```



image 1/1 /content/car-plate-detection/images/Cars181.png: 320x192 1 license_plate, 46.9ms
Speed: 0.9ms preprocess, 46.9ms inference, 0.9ms postprocess per image at shape (1, 3, 320, 192)
Detected text: SHIT HAD



```
predict_and_plot(test.iloc[25].img_path)
```



image 1/1 /content/car-plate-detection/images/Cars198.png: 256x320 1 license_plate, 54.7ms
Speed: 1.1ms preprocess, 54.7ms inference, 0.9ms postprocess per image at shape (1, 3, 256, 320)
Detected text: ~-MH01AV8866)



✓ Conclusion:

This experiment demonstrates how **cognitive image-based applications** can enhance the **Apartment Assistant** system. By combining **image classification, OCR, and license plate detection**, the system can acquire knowledge from real-world apartment-related images. The addition of **vehicle plate detection** improves apartment security and makes parking and visitor management more efficient. Such applications show the power of cognitive computing in **real-life problem-solving** for smarter living spaces.